

TOP-K IS NOT A BUDGET FOR HYBRID RETRIEVAL

Chunran Zhang

School of Computing and Artificial Intelligence
Southwest Jiaotong University, Chengdu, China
chronis@my.swjtu.edu.cn

ABSTRACT

Hybrid retrieval for RAG often fuses fixed-depth dense and sparse rankings. Exact fusion avoids this truncation, but completing a fixed Top- K can require sharply different numbers of accesses across queries. DiBud (Direct Budgeting) takes an access budget directly and returns the certified prefix available within it. Incremental certification determines the output size, while selective access resolves remaining ranking uncertainty. Replays on five query sets show that, at 2048 accesses, DiBud certifies 7.86% more results within the first 100 positions than balanced access. After budget calibration for 95% of exact Top-20 quality, held-out NFCorpus and SciFact queries retain 95.81% and 97.68% of nDCG@20 while reducing accesses by 65.92% and 83.41%. Five corpus snapshots and observation-limited cases expose variation in output size and the limits of this tradeoff.

Index Terms— Hybrid retrieval, reciprocal rank fusion, budgeted retrieval, exact prefix, RAG

1. INTRODUCTION

Hybrid retrieval for retrieval-augmented generation (RAG) commonly combines the top L results from dense and sparse retrievers using reciprocal rank fusion (RRF) [1, 2]. A fixed window limits the number of entries accessed but also excludes contributions beyond that window. When a document appears in one ranked list but not in the other list’s window, fusion assigns zero to the latter contribution. Yet absence from the window only means that the rank has not been observed; its contribution remains unknown. Replacing an unknown contribution with zero can change the output order even when the candidate set already contains every document in the top K under RRF over the full ranked lists. Thus, L is a retrieval-depth parameter that affects both retrieval quality and access cost.

In practice, L can be selected using historical queries to balance quality and cost. However, RAG systems serving open-ended requests over evolving knowledge bases continually encounter new queries and updated corpora. Both changes alter the ranked lists and the positions of cross-channel contributions, making it difficult for a previously se-

lected L to preserve the original quality–cost tradeoff. Exact rank aggregation maintains bounds on unobserved contributions and accesses the ranked lists until the required results are determined [3]. EAHR applies this approach to hybrid retrieval to obtain the exact top K in order without a predefined truncation depth [4]. Nevertheless, the number of accesses required can vary sharply across queries: resolving competition at just a few positions can require accessing entries far down the ranked lists.

Because the same K can require substantially different numbers of accesses, resource-constrained requests need to bound accesses directly rather than control cost indirectly through K . We take an access budget as input, defined as the maximum total number of sorted accesses to the two ranked lists, without specifying the output size in advance. Each sorted access retrieves the next entry from one list. The system incrementally certifies results and returns a certified prefix when the budget is exhausted: an initial segment guaranteed to match the RRF ranking over the full lists in both membership and order. Differences in query difficulty therefore translate into different output sizes, rather than additional accesses required to complete a fixed top- K result.

To support retrieval under an access budget, we introduce DiBud (Direct Budgeting). DiBud incrementally computes the fused ranking, appending results whose ranks have been certified [5] and directing subsequent accesses toward unresolved competition at later positions [6]. When the budget is exhausted, it returns the current certified prefix. No K is specified in advance; the output size is determined by the ranking information obtained within the budget.

We evaluate this approach through replays on five query sets and five corpus snapshots. We first examine fixed-depth transfer, then measure the access cost and incremental certification of exact top- K results. Under equal budgets, we compare DiBud with balanced sorted access, including output distributions and changes across snapshots. Finally, calibration on disjoint queries quantifies the relevance retained and accesses saved relative to completing exact Top-20.

2. DIBUD: BUDGETED EXACT-PREFIX FUSION

Given a query, a corpus snapshot, and a read budget B , DiBud reads the dense and sparse rankings on demand and returns a certified prefix of their complete RRF ranking. The budget determines how much can be read; certification determines how much can be returned.

2.1. Budget and output

Let $r_i(x)$ denote the one-based rank of document x in channel $i \in \{D, S\}$, and let $g(r) = 1/(c+r)$ for rank constant $c \geq 0$. Set $r_i(x) = \infty$ and $g(\infty) = 0$ if x is outside that channel's support. The complete fusion score is

$$F(x) = \sum_{i \in \{D, S\}} g(r_i(x)). \quad (1)$$

Sorting the union of the channel supports by decreasing $F(x)$, with a fixed document-identifier order to break ties, defines the complete sequence T^* .

Both channels supply resumable exact ranking prefixes from the same snapshot. If d_D and d_S are their read depths and A is the output sequence, the required contract is

$$d_D + d_S \leq B, \quad A = T_{1:K}^*, \quad K = |A|. \quad (2)$$

K is not specified in advance. The output is empty if no position can be certified. The budget counts channel entries read, including separate reads of the same document in different channels.

2.2. Prefix certification

Let S_i be the documents already read from channel i . An unread contribution in that channel is bounded by

$$u_i = \begin{cases} 0, & \text{if confirmed exhausted,} \\ g(d_i + 1), & \text{otherwise.} \end{cases} \quad (3)$$

For an observed document x , the score bounds are

$$\ell(x) = \sum_{i: x \in S_i} g(r_i(x)), \quad (4)$$

$$h(x) = \ell(x) + \sum_{i: x \notin S_i} u_i. \quad (5)$$

A document unseen in both channels has score at most $h_\emptyset = u_D + u_S$.

Among observed documents not yet emitted, select x with the highest lower bound, breaking ties by the fixed identifier order. Certify x as the next output when its lower bound exceeds every remaining competitor's upper bound and $h_\emptyset$. Equality with an observed competitor is allowed only when x wins the identifier tie; comparison with the unseen bound

remains strict. Append x and repeat until the next position cannot be certified.

The bounds ensure $\ell(x) \leq F(x) \leq h(x)$. Each emitted document therefore precedes every remaining document, including any not yet observed. Applying this argument at each output position establishes the prefix equality in (2).

2.3. Reading and stopping

When the next position cannot be certified, let y be the observed competitor with the highest upper bound, excluding x and the emitted prefix and breaking ties by identifier. If $h(y) \geq h_\emptyset$ and y is missing a contribution from exactly one channel, advance that channel. This either reveals the contribution or tightens its upper bound.

Otherwise, advance the channel whose next batch offers the larger reduction in the unread bound. For a batch of b entries in a channel that remains open after the read, this reduction is

$$\Delta_i = g(d_i + 1) - g(d_i + b + 1). \quad (6)$$

If exhaustion is established, the new bound is zero. Exhausted channels are skipped, and equal reductions are resolved by a fixed channel order (dense first). This rule also initializes reading when no document has been observed.

After each batch, update the bounds and extend the certified prefix. The batch size is capped by the remaining budget. When the budget is consumed, or both channels are exhausted, return the accumulated output. Certification depends on the observed bounds, so changing the read schedule affects the number of certified results without changing the meaning of an emitted position.

3. EXPERIMENTS

3.1. Setup

We evaluate 770 queries from five query sets—TREC-DL 2019/2020, NFCorpus, SciFact, and TREC-COVID—and the same 30 queries across five TREC-COVID snapshots [7, 8, 9, 10, 11]. We replay exact rankings from bge-small-en-v1.5 [12] and BM25 [13] ($k_1 = 0.9$, $b = 0.4$), using equal-weight RRF with contributions $1/(59 + r)$ and deterministic tie breaking.

First, we choose L on calibration queries and test it on held-out queries and later corpus snapshots. Next, we measure how many accesses are needed to complete exact Top- k and when each result becomes certain. We then compare DiBud with Balanced under the same budgets and certification rules. Balanced alternates between the two lists. Budgets range from 128 to 10000, and we count certified results up to 20 and 100. Both methods access one entry at a time; we also test batches of 64.

Table 1. Sorted accesses to exact Top-20 under unit SNRA. Completion is observed within the saved rankings and 10,000-access limit; unresolved P95 values are lower bounds.

Query set	n	Median	P95	Complete (%)
TREC-DL 2019	43	311	≥ 9108	81.4
TREC-DL 2020	54	444	≥ 9912	85.2
NFCorpus	323	308	2685	100.0
SciFact	300	325	3362	100.0
TREC-COVID	50	2544	≥ 9842	66.0

Finally, we compare budgeted stopping with completing exact Top-20. Budgets are selected on calibration queries to retain 90%, 95%, or 99% of the reference mean $nDCG@20$. We report the quality retained and accesses saved on held-out queries.

Results are averaged within each query set, then equally across sets. Where reported, 95% confidence intervals use 2000 query bootstrap samples. Queries that reach the end of the saved rankings remain in the analysis, with bounds reported for outcomes that cannot be determined.

3.2. Transfer of fixed depths

We evaluate $L \in \{10, 20, 50, 100, 200, 500, 1000, 2000, 5000\}$. On the same MS MARCO corpus, increasing L from 20 to 5000 raises $nDCG@10$ from 0.6538 to 0.6812 for TREC-DL 2019, but lowers it from 0.6424 to 0.6243 for TREC-DL 2020 (Fig. 1a). Changing queries reverses the benefit of a deeper window.

For the same 30 queries across five TREC-COVID snapshots, the best tested depths are 50, 100, 200, 200, and 5000 (Fig. 1b). We then test whether an earlier selection transfers: five-fold calibration selects L using Round-1 queries and freezes it for held-out queries across all rounds. The $nDCG@10$ difference from full RRF changes from -0.0085 in Round 1 to -0.0272 in Round 5 (95% CI: $[-0.0699, -0.0058]$; Fig. 1c). A depth selected on earlier data does not necessarily maintain its relative effectiveness as the corpus changes.

3.3. Fixed-K cost and incremental certification

For exact Top-20, median access counts range from 308 to 2544 (Table 1). All NFCorpus and SciFact queries complete within the observed range, but their P95 costs are 2685 and 3362, respectively: 8.72 and 10.34 times their medians. On the two TREC-DL sets and TREC-COVID, 18.60%, 14.81%, and 34.00% of queries do not complete within the observed range. Their costs cannot be replaced by statistics over completed queries. Among queries with observed Top-20 completion, half of the eventual accesses already certify 11.87 results on an equal-set average. The largest single certification increment accounts for 25.42% of total accesses on average.

Table 2. Mean certified output at $B = 2048$. Both methods use unit accesses and the same certifier. All endpoints in this table are observed.

Query set	Cap 100		Cap 20	
	Balanced	DiBud	Balanced	DiBud
TREC-DL 2019	41.86	46.19	17.77	17.91
TREC-DL 2020	34.91	38.31	18.33	18.35
NFCorpus	62.06	63.71	19.24	19.28
SciFact	43.55	47.26	19.44	19.57
TREC-COVID	26.48	29.80	15.58	15.78

Thus, useful prefixes are available before the fixed-size result is complete, although the measurements do not imply that every query is dominated by one extreme position (Fig. 2).

3.4. Certified output under equal budgets

At $B = 2048$, DiBud increases mean certified output within the first 100 positions from 41.77 to 45.05, a 7.86% gain over Balanced (Table 2). Within the first 20 positions, the mean changes from 18.07 to 18.18. All five sets gain output, with larger improvements for longer prefixes. The macro fraction returning at least 20 results increases from 75.10% to 76.22%; average prefix growth and complete-20 delivery are distinct outcomes. Over the full log-budget range, the lower bound on DiBud’s normalized certified-prefix area exceeds Balanced’s upper bound on every static set. Some asymmetric paths reach the saved-list limit at large budgets, so Fig. 3 shows intervals rather than treating the retained prefix as the final output. The lower bound on the area difference remains positive when both schedules use batches of 64. The gain is therefore not solely a unit-access effect.

Across the five snapshots at $B = 2048$, DiBud returns means of 18.00, 17.57, 16.40, 15.77, and 15.10 results within the first 20 positions, versus 17.70, 17.20, 16.03, 15.30, and 14.83 for Balanced. A fixed budget yields different output sizes as the corpus changes. DiBud improves the mean in every round. Full budget curves, output quantiles, empty-output rates, and delivery thresholds are retained with the experiment artifacts.

3.5. Quality and access cost relative to exact Top-20

We compare two stopping conditions on the same unit-SNRA trajectory: complete exact Top-20, or stop at the budget or at 20 certified results, whichever comes first. Evaluation uses linear-gain $nDCG@20$, with absent positions contributing zero. For each query set, five-fold calibration chooses the smallest budget retaining 90%, 95%, or 99% of the reference mean $nDCG@20$; the frozen choice is evaluated on held-out queries. Retention is a ratio of means, not a mean of per-query ratios with an arbitrary value for zero-reference queries. At the 95% calibration target, NFCorpus and Sci-




Fig. 1. Depth sensitivity and transfer. (a) Different queries on the same corpus. (b) The same queries across corpus snapshots, evaluated with each round’s judgments. Circled points mark retrospectively selected best depths. (c) Round-1 depths applied to held-out queries across rounds. Error bars show nested bootstrap 95% confidence intervals, including depth reselection and preserving query identities across rounds.




Fig. 2. Exact Top-20 completion and earlier certified output. (a) Observed completions divided by all queries; plateaus retain unresolved queries in the denominator. (b) Output at fractions of each query’s observed completion cost, conditional on completion. This retrospective normalization is not an on-line policy.

Table 3. Held-out Top-20 comparison after selecting the smallest budget retaining 95% of reference mean nDCG@20 on calibration queries. Retention is a ratio of dataset means. Intervals and $\geq$ reflect unobserved continuations, not confidence intervals.

Query set	Retained (%)	Fewer accesses (%)	Top20 (%)
TREC-DL 2019	95.05	≥ 76.02	67.4
TREC-DL 2020	95.89	≥ 70.96	68.5
NFCorpus	95.81	65.92	52.3
SciFact	97.68	83.41	11.3
TREC-COVID	95.14–96.43	≥ 5.30	64.0

Fact retain 95.81% and 97.68% of held-out nDCG@20 while reducing accesses by 65.92% and 83.41% (Table 3). TREC-DL 2019/2020 retain 95.05% and 95.89%. Their all-query access reductions are at least 76.02% and 70.96%, using lower bounds on unresolved exact Top-20 costs. Within the completed-query subsets alone, reductions are 35.70% and 47.07%, illustrating the influence of difficult queries on total cost. TREC-COVID requires a different tradeoff: retention is bounded by 95.14%–96.43%, while access reduction is guaranteed only to be at least 5.30%. At the 99% target, five, four, and five calibration folds on TREC-DL 2019, TREC-DL 2020, and TREC-COVID fail to establish the target from the available rankings. These are not successful target crossings.

Budgeted stopping can therefore retain much of the top-ranked relevance gain while reducing accesses spent completing Top-20, at the cost of output size. At the 95% target, only 52.32% of NFCorpus queries and 11.33% of SciFact queries return all 20 results. Quality retention is not a substitute for




Fig. 3. Certified output versus budget and snapshot. Shading bounds output after a saved-list observation limit; it is not a confidence interval. Static panels use a reporting cap of 100. The snapshot panel uses cap 20 and $B = 2048$, where all endpoints are observed.

result coverage or a guarantee for each query.

4. CONCLUSION

Fixed-depth fusion can change its relative effectiveness across queries and corpus snapshots, while exact Top- K completion retains a variable access cost. DiBud directly bounds sorted accesses and lets incremental certification determine the returned prefix. Matched-budget replays show increased output over balanced access, and held-out budget calibration demonstrates substantial access savings at high mean quality retention on several query sets. The tradeoff varies across collections: fewer accesses can mean fewer returned results, and a quality target need not transfer to every query. The evidence concerns logical accesses to exact ranked lists; online retrieval latency remains outside this evaluation.

5. ACKNOWLEDGMENTS

OpenAI Codex assisted with drafting, English translation, and language editing of the abstract and Sections 1–4 from the research notes and Chinese manuscript.




Fig. 4. Held-out quality and access reduction relative to exact Top-20 after 95% calibration. The dashed line is the calibration target, not a per-query guarantee. A range denotes unobserved continuation; reduction labels with $\geq$ are conservative bounds.

6. REFERENCES

- [1] Elastic, “Reciprocal rank fusion,” <https://www.elastic.co/docs/reference/elasticsearch/rest-apis/reciprocal-rank-fusion>, Accessed September 10, 2026.
- [2] Gordon V. Cormack, Charles L. A. Clarke, and Stefan Buettcher, “Reciprocal rank fusion outperforms condorcet and individual rank learning methods,” in *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2009, pp. 758–759.
- [3] Ronald Fagin, Amnon Lotem, and Moni Naor, “Optimal aggregation algorithms for middleware,” *Journal of Computer and System Sciences*, vol. 66, no. 4, pp. 614–656, 2003.
- [4] Chunran Zhang, “Exact adaptive hybrid retrieval without fixed top-L cutoffs,” *arXiv preprint arXiv:2608.07152*, 2026.
- [5] Nikos Mamoulis, Kit Hung Cheng, Man Lung Yiu, and David W. Cheung, “Efficient aggregation of ranked inputs,” in *Proc. IEEE International Conference on Data Engineering (ICDE)*, 2006, p. 72.
- [6] Jing Yuan, Guang-Zhong Sun, Ye Tian, Guoliang Chen, and Zhi Liu, “Selective-NRA algorithms for top- k queries,” in *Advances in Data and Web Management (APWeb/WAIM)*, 2009, pp. 15–26.
- [7] Nick Craswell, Bhaskar Mitra, Emine Yilmaz, Daniel Campos, and Ellen M. Voorhees, “Overview of the TREC 2019 deep learning track,” *arXiv preprint arXiv:2003.07820*, 2020.
- [8] Nick Craswell, Bhaskar Mitra, Emine Yilmaz, and Daniel Campos, “Overview of the TREC 2020 deep learning track,” *arXiv preprint arXiv:2102.07662*, 2021.
- [9] Vera Boteva, Demian Gholipour, Artem Sokolov, and Stefan Riezler, “A full-text learning to rank dataset for medical information retrieval,” in *Advances in Information Retrieval*. 2016, pp. 716–722, Springer.
- [10] David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hannaneh Hajishirzi, “Fact or fiction: Verifying scientific claims,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*. 2020, pp. 7534–7550, Association for Computational Linguistics.
- [11] Kirk Roberts, Tasmeer Alam, Steven Bedrick, Dina Demner-Fushman, Kyle Lo, Ian Soboroff, Ellen Voorhees, Lucy Lu Wang, and William R. Hersh, “Searching for scientific evidence in a pandemic: An overview of TREC-COVID,” *Journal of Biomedical Informatics*, vol. 121, pp. 103865, 2021.
- [12] BAAI, “bge-small-en-v1.5: Model card,” <https://huggingface.co/BAAI/bge-small-en-v1.5>, Accessed September 10, 2026.
- [13] Stephen Robertson and Hugo Zaragoza, “The probabilistic relevance framework: BM25 and beyond,” *Foundations and Trends in Information Retrieval*, vol. 3, no. 4, pp. 333–389, 2009.